

PRODUCT DEMAND PREDICTION WITH **MACHINE LEARNINGS:**

912421104002 : ABIRAMI.M

PHASE-2 PROJECT SUBMISSION DOCUMENT

PROJECT : PRODUCT DEMAND PERDICTION



INTRODUCTION:

- Product demand prediction with machine learning involves using algorithms and statistical models to forecast the future demand for a specific product.
- This is a valuable process for businesses in various industries, as it helps them optimize inventory management, production planning, and supply chain operations.
- Machine learning techniques can analyze historical data and various factors influencing demand to make accurate predictions.
- Product demand prediction with machine learning can help businesses reduce excess inventory costs, minimize stockouts, improve customer satisfaction, and enhance overall operational efficiency.

CONTENT FOR PHASE-2 PROJECT:

Consider incorporating time series forecasting techniques like ARIMA or Prophet to capture temporal patterns in demand data.

DATA SOURCE:

A good data source for product demand prediction using machine learning is:

Dataset Link:

<https://www.kaggle.com/datasets/chakradharmattapalli/product-demand-prediction-with-machine-learning>

| ID | Store ID | Total Price | Base Price | Units Sold |
|----|----------|-------------|------------|------------|
| 1 | 8091 | 99.0375 | 111.8625 | 20 |
| 2 | 8091 | 99.0375 | 99.0375 | 28 |
| 3 | 8091 | 133.95 | 133.95 | 19 |
| 4 | 8091 | 133.95 | 133.95 | 44 |
| 5 | 8091 | 141.075 | 141.075 | 52 |
| 9 | 8091 | 227.2875 | 227.2875 | 18 |
| 10 | 8091 | 327.0375 | 327.0375 | 47 |
| 13 | 8091 | 210.9 | 210.9 | 50 |
| 14 | 8091 | 190.2375 | 234.4125 | 82 |
| 17 | 8095 | 99.0375 | 99.0375 | 99 |
| 18 | 8095 | 97.6125 | 97.6125 | 120 |
| 19 | 8095 | 98.325 | 98.325 | 40 |
| 22 | 8095 | 133.2375 | 133.2375 | 68 |
| 23 | 8095 | 133.95 | 133.95 | 87 |
| 24 | 8095 | 139.65 | 139.65 | 186 |
| 27 | 8095 | 236.55 | 280.0125 | 54 |
| 28 | 8095 | 214.4625 | 214.4625 | 74 |
| 29 | 8095 | 266.475 | 296.4 | 102 |
| 30 | 8095 | 173.85 | 192.375 | 214 |
| 31 | 8095 | 205.9125 | 205.9125 | 28 |
| 32 | 8095 | 205.9125 | 205.9125 | 7 |
| 33 | 8095 | 248.6625 | 248.6625 | 48 |
| 34 | 8095 | 200.925 | 200.925 | 78 |
| 35 | 8095 | 190.2375 | 240.825 | 57 |
| 37 | 8095 | 427.5 | 448.1625 | 50 |
| 38 | 8095 | 429.6375 | 458.1375 | 62 |
| 39 | 8095 | 177.4125 | 177.4125 | 22 |
| 42 | 8094 | 87.6375 | 87.6375 | 109 |
| 43 | 8094 | 88.35 | 88.35 | 133 |
| 44 | 8094 | 85.5 | 85.5 | 11 |
| 45 | 8094 | 128.25 | 180.975 | 9 |
| 47 | 8094 | 127.5375 | 127.5375 | 19 |
| 48 | 8094 | 123.975 | 123.975 | 33 |
| 49 | 8094 | 139.65 | 164.5875 | 49 |
| 50 | 8094 | 235.8375 | 235.8375 | 32 |
| 51 | 8094 | 234.4125 | 234.4125 | 47 |
| 52 | 8094 | 235.125 | 235.125 | 27 |
| 53 | 8094 | 227.2875 | 227.2875 | 69 |
| 54 | 8094 | 312.7875 | 312.7875 | 49 |

EXPLANATION FOR PHASE-2 PROJECT:

Data Collection and Preprocessing:

Gather historical demand data, ensuring that it is time-stamped and organized chronologically. Preprocess the data by addressing missing values, outliers, and any other data quality issues.

Exploratory Data Analysis (EDA):

Conduct EDA to understand the temporal patterns and characteristics of the demand data. Look for seasonality, trends, and other recurring patterns. Visualization tools and statistical tests can be helpful in this phase.

Incorporating time series forecasting techniques:

➤ ARIMA (Auto Regressive Integrated Moving Average):

Suitable for stationary data with autoregressive and moving average components.

➤ SARIMA (Seasonal ARIMA):

Extends ARIMA to handle seasonal patterns in data.

➤ Exponential Smoothing Methods:

These include Holt-Winters for capturing trends and seasonality.

➤ **Prophet:**

Developed by Facebook, Prophet is useful for data with daily observations, holidays, and seasonality.

➤ **Deep Learning Models (e.g., LSTM and GRU):**

Suitable for capturing complex temporal patterns, but they may require more data and computational resources.

Model Training:

Train the selected time series forecasting model using historical demand data. This involves estimating model parameters and seasonal components, if applicable.

Validation and Hyperparameter Tuning:

Assess the model's performance using validation data or cross-validation. Fine-tune hyperparameters and adjust the model structure as needed to improve forecasting accuracy.

Forecasting:

Once the model is trained and validated, use it to make predictions for future time periods. These forecasts will capture temporal patterns and provide insights into expected demand behavior.

Performance Evaluation:

Evaluate the forecasting model's performance using appropriate metrics like Mean Absolute Error (MAE), Mean Squared Error (MSE), Root Mean Squared Error (RMSE), and forecast accuracy measures.

Continuous Monitoring and Updating:

Implement a process for regularly updating and retraining the model as new demand data becomes available. This ensures that the model adapts to changing demand patterns over time.

Incorporate External Factors:

Consider adding external variables such as promotional activities, economic indicators, or weather data to your model to account for factors that influence demand fluctuations.

PROGRAM:

Product Demand Prediction:

```
import pandas as pd
import numpy as np
import plotly.express as px
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
```

```
data=pd.read_csv("C:\Users\mabir\AppData\Local\Microsoft\Windows\INetCache\IE\AHLGJQP8\archive[1].zip ")
```

```
data.head()
```

Output:

| ID | Store ID | Total Price | Base Price | Units Sold |
|----|----------|-------------|------------|------------|
| 1 | 8091 | 99.0375 | 111.8625 | 20 |
| 2 | 8091 | 99.0375 | 99.0375 | 28 |
| 3 | 8091 | 133.95 | 133.95 | 19 |
| 4 | 8091 | 133.95 | 133.95 | 44 |
| 5 | 8091 | 141.075 | 141.075 | 52 |
| 9 | 8091 | 227.2875 | 227.2875 | 18 |
| 10 | 8091 | 327.0375 | 327.0375 | 47 |
| 13 | 8091 | 210.9 | 210.9 | 50 |
| 14 | 8091 | 190.2375 | 234.4125 | 82 |
| 17 | 8095 | 99.0375 | 99.0375 | 99 |
| 18 | 8095 | 97.6125 | 97.6125 | 120 |
| 19 | 8095 | 98.325 | 98.325 | 40 |
| 22 | 8095 | 133.2375 | 133.2375 | 68 |
| 23 | 8095 | 133.95 | 133.95 | 87 |
| 24 | 8095 | 139.65 | 139.65 | 186 |
| 27 | 8095 | 236.55 | 280.0125 | 54 |
| 28 | 8095 | 214.4625 | 214.4625 | 74 |
| 29 | 8095 | 266.475 | 296.4 | 102 |
| 30 | 8095 | 173.85 | 192.375 | 214 |
| 31 | 8095 | 205.9125 | 205.9125 | 28 |
| 32 | 8095 | 205.9125 | 205.9125 | 7 |
| 33 | 8095 | 248.6625 | 248.6625 | 48 |
| 34 | 8095 | 200.925 | 200.925 | 78 |
| 35 | 8095 | 190.2375 | 240.825 | 57 |
| 37 | 8095 | 427.5 | 448.1625 | 50 |
| 38 | 8095 | 429.6375 | 458.1375 | 62 |
| 39 | 8095 | 177.4125 | 177.4125 | 22 |
| 42 | 8094 | 87.6375 | 87.6375 | 109 |
| 43 | 8094 | 88.35 | 88.35 | 133 |
| 44 | 8094 | 85.5 | 85.5 | 11 |
| 45 | 8094 | 128.25 | 180.975 | 9 |
| 47 | 8094 | 127.5375 | 127.5375 | 19 |
| 48 | 8094 | 123.975 | 123.975 | 33 |
| 49 | 8094 | 139.65 | 164.5875 | 49 |
| 50 | 8094 | 235.8375 | 235.8375 | 32 |


| | | | | |
|----|------|----------|----------|----|
| 51 | 8094 | 234.4125 | 234.4125 | 47 |
| 52 | 8094 | 235.125 | 235.125 | 27 |
| 53 | 8094 | 227.2875 | 227.2875 | 69 |
| 54 | 8094 | 312.7875 | 312.7875 | 49 |

Relationship between price and demand for the product:

```
fig = px.scatter(data, x="Units Sold", y="Total Price",
                size='Units Sold')
```

```
fig.show()
```

output:



Correlation between the features of the dataset:

```
print(data.corr())
```

Output:

```

      ID  Store ID  Total Price  Base Price  Units Sold
ID      1.000000  0.007464    0.008473   0.018932  -
0.010616
Store ID  0.007464  1.000000   -0.038315  -0.038848  -
0.004372
```



```

Total Price 0.008473 -0.038315 1.000000 0.958885 -
0.235625
Base Price 0.018932 -0.038848 0.958885 1.000000 -
0.140032
Units Sold -0.010616 -0.004372 -0.235625 -0.140032
1.000000

```

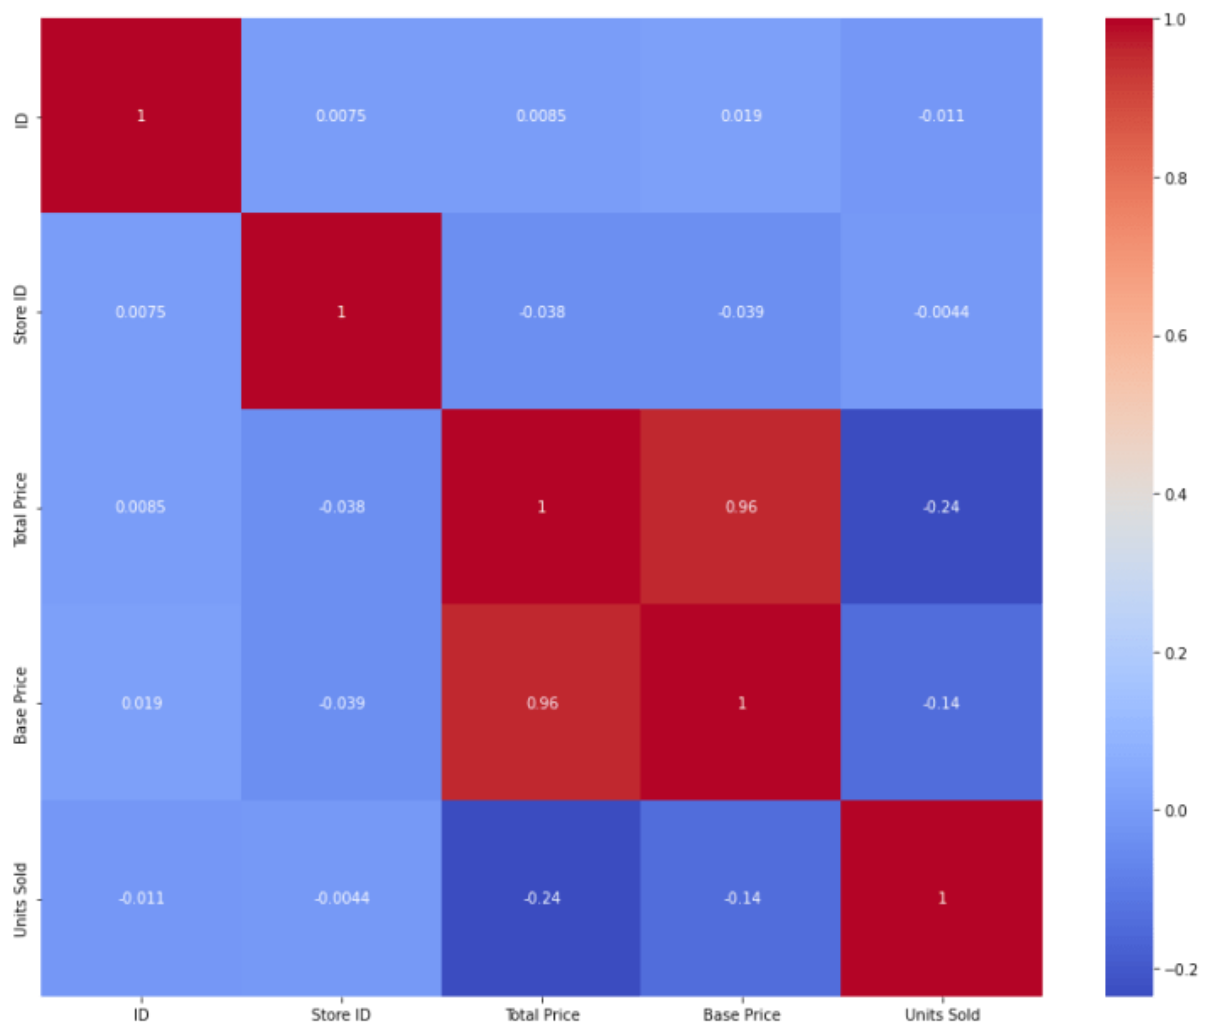
```
correlations = data.corr(method='pearson')
```

```
plt.figure(figsize=(15, 12))
```

```
sns.heatmap(correlations, cmap="coolwarm", annot=True)
```

```
plt.show()
```

Output:



```
# fit an ARIMA model and plot residual errors

from pandas import datetime
from pandas import read_csv
from pandas import DataFrame
from statsmodels.tsa.arima.model import ARIMA
from matplotlib import pyplot
# load dataset

def parser(x):
    return datetime.strptime('190'+x, '%Y-%m')
series = read_csv('shampoo-sales.csv', header=0, index_col=0,
parse_dates=True, squeeze=True, date_parser=parser)
series.index = series.index.to_period('M')
# fit model

model = ARIMA(series, order=(5,1,0))
model_fit = model.fit()
# summary of fit model

print(model_fit.summary())
# line plot of residuals

residuals = DataFrame(model_fit.resid)
residuals.plot()
pyplot.show()
# density plot of residuals

residuals.plot(kind='kde')
pyplot.show()
# summary stats of residuals

print(residuals.describe())
```

Output:

SARIMAX Results

```

=====
=====
Dep. Variable:          Sales  No. Observations:          36
Model:                ARIMA(5, 1, 0)  Log Likelihood      -198.485
Date:                Thu, 10 Dec 2020  AIC                408.969
Time:                09:15:01  BIC                418.301
Sample:                01-31-1901  HQIC                412.191
                        - 12-31-1903
Covariance Type:      opg
=====
=====

```


| | coef | std err | z | P> z | [0.025 | 0.975] |
|--------|-----------|----------|--------|-------|----------|----------|
| ar.L1 | -0.9014 | 0.247 | -3.647 | 0.000 | -1.386 | -0.417 |
| ar.L2 | -0.2284 | 0.268 | -0.851 | 0.395 | -0.754 | 0.298 |
| ar.L3 | 0.0747 | 0.291 | 0.256 | 0.798 | -0.497 | 0.646 |
| ar.L4 | 0.2519 | 0.340 | 0.742 | 0.458 | -0.414 | 0.918 |
| ar.L5 | 0.3344 | 0.210 | 1.593 | 0.111 | -0.077 | 0.746 |
| sigma2 | 4728.9608 | 1316.021 | 3.593 | 0.000 | 2149.607 | 7308.314 |

```


=====
=====
Ljung-Box (L1) (Q):          0.61  Jarque-Bera (JB):          0.96
Prob(Q):                    0.44  Prob(JB):                  0.62
Heteroskedasticity (H):      1.07  Skew:                  0.28
Prob(H) (two-sided):         0.90  Kurtosis:              2.41
=====
=====

```

First, we get a line plot of the residual errors, suggesting that there may still be some trend information not captured by the model.



Next, we get a density plot of the residual error values, suggesting the errors are Gaussian, but may not be centered on zero.



Rolling Forecast ARIMA Model:

evaluate an ARIMA model using a walk-forward validation

```
from pandas import read_csv
from pandas import datetime
from matplotlib import pyplot
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error
from math import sqrt

# load dataset

def parser(x):
    return datetime.strptime('190'+x, '%Y-%m')
```

```
series = read_csv('shampoo-sales.csv', header=0, index_col=0,
parse_dates=True, squeeze=True, date_parser=parser)
series.index = series.index.to_period('M')
# split into train and test sets

X = series.values
size = int(len(X) * 0.66)
train, test = X[0:size], X[size:len(X)]
history = [x for x in train]
predictions = list()
# walk-forward validation

for t in range(len(test)):
    model = ARIMA(history, order=(5,1,0))
    model_fit = model.fit()
    output = model_fit.forecast()
    yhat = output[0]
    predictions.append(yhat)
    obs = test[t]
    history.append(obs)
    print('predicted=%f, expected=%f' % (yhat, obs))
# evaluate forecasts

rmse = sqrt(mean_squared_error(test, predictions))
print('Test RMSE: %.3f' % rmse)
# plot forecasts against actual outcomes
```

```
pyplot.plot(test)
pyplot.plot(predictions, color='red')
pyplot.show()
```

Running the example prints the prediction and expected value each iteration.

We can also calculate a final root mean squared error score (RMSE) for the predictions, providing a point of comparison for other ARIMA configurations.

```
predicted=343.272180, expected=342.300000
```

```
predicted=293.329674, expected=339.700000
```

```
predicted=368.668956, expected=440.400000
```

```
predicted=335.044741, expected=315.900000
```

```
predicted=363.220221, expected=439.300000
```

```
predicted=357.645324, expected=401.300000
```

```
predicted=443.047835, expected=437.400000
```

```
predicted=378.365674, expected=575.500000
```

```
predicted=459.415021, expected=407.600000
```

```
predicted=526.890876, expected=682.000000
```


```
predicted=457.231275, expected=475.300000
```

```
predicted=672.914944, expected=581.300000
```

```
predicted=531.541449, expected=646.900000
```

```
Test RMSE: 89.021
```

A line plot is created showing the expected values (blue) compared to the rolling forecast predictions (red). We can see the values show some trend and are in the correct scale.



Seasonal ARIMA (SARIMA):


```
SARIMAXmodel = SARIMAX(y, order = (5, 4, 2),  
seasonal_order=(2,2,2,12))  
SARIMAXmodel = SARIMAXmodel.fit()
```

```
y_pred = SARIMAXmodel.get_forecast(len(test.index))  
y_pred_df = y_pred.conf_int(alpha = 0.05)  
y_pred_df["Predictions"] = SARIMAXmodel.predict(start =  
y_pred_df.index[0], end = y_pred_df.index[-1])  
y_pred_df.index = test.index  
y_pred_out = y_pred_df["Predictions"]  
plt.plot(y_pred_out, color='Blue', label = 'SARIMA Predictions')
```



```
plt.legend()
```

Output:



Prophet:

make an in-sample forecast

```
from pandas import read_csv
```

```
from pandas import to_datetime
```

```
from pandas import DataFrame
```

```
from fbprophet import Prophet
```

```
from matplotlib import pyplot
```

load data

```
path =  
'https://raw.githubusercontent.com/jbrownlee/Datasets/master/monthly-car-sales.csv'  
  
df = read_csv(path, header=0)  
  
# prepare expected column names  
  
df.columns = ['ds', 'y']  
df['ds']= to_datetime(df['ds'])  
  
# define the model  
  
model = Prophet()  
  
# fit the model  
  
model.fit(df)  
  
# define the period for which we want a prediction  
  
future = list()  
for i in range(1, 13):  
    date = '1968-%02d' % i  
    future.append([date])  
future = DataFrame(future)  
future.columns = ['ds']  
future['ds']= to_datetime(future['ds'])  
  
# use the model to make a forecast
```

```
forecast = model.predict(future)
```

```
# summarize the forecast
```

```
print(forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']].head())
```

```
# plot forecast
```


```
model.plot(forecast)
```

```
pyplot.show()
```

Running the example forecasts the last 12 months of the dataset.

The first five months of the prediction are reported and we can see that values are not too different from the actual sales values in the dataset(output).

| | ds | yhat | yhat_lower | yhat_upper |
|---|------------|--------------|--------------|--------------|
| 0 | 1968-01-01 | 14364.866157 | 12816.266184 | 15956.555409 |
| 1 | 1968-02-01 | 14940.687225 | 13299.473640 | 16463.811658 |
| 2 | 1968-03-01 | 20858.282598 | 19439.403787 | 22345.747821 |
| 3 | 1968-04-01 | 22893.610396 | 21417.399440 | 24454.642588 |
| 4 | 1968-05-01 | 24212.079727 | 22667.146433 | 25816.191457 |



Tying this together, the example below demonstrates how to evaluate a Prophet model on a hold-out dataset.

```
# evaluate prophet time series forecasting model on hold out dataset
```

```
from pandas import read_csv
```

```
from pandas import to_datetime
```

```
from pandas import DataFrame
```

```
from fbprophet import Prophet
```

```
from sklearn.metrics import mean_absolute_error
```

```
from matplotlib import pyplot
```

```
# load data
```

```
path =
```

```
'https://raw.githubusercontent.com/jbrownlee/Datasets/master/monthly-car-sales.csv'
```

```
df = read_csv(path, header=0)

# prepare expected column names
df.columns = ['ds', 'y']

df['ds'] = to_datetime(df['ds'])

# create test dataset, remove last 12 months
train = df.drop(df.index[-12:])

print(train.tail())

# define the model
model = Prophet()

# fit the model
model.fit(train)

# define the period for which we want a prediction
future = list()

for i in range(1, 13):
    date = '1968-%02d' % i
    future.append([date])

future = DataFrame(future)

future.columns = ['ds']

future['ds'] = to_datetime(future['ds'])

# use the model to make a forecast
forecast = model.predict(future)
```

```
# calculate MAE between expected and predicted values for  
december
```

```
y_true = df['y'][-12:].values
```

```
y_pred = forecast['yhat'].values
```

```
mae = mean_absolute_error(y_true, y_pred)
```

```
print('MAE: %.3f' % mae)
```

```
# plot expected vs actual
```


```
pyplot.plot(y_true, label='Actual')
```

```
pyplot.plot(y_pred, label='Predicted')
```

```
pyplot.legend()
```

```
pyplot.show()
```

Output:



CONCLUSION AND FUTURE WORK(PHASE-2):

- In this phase-2 project conclusion, we will summarize the key findings and insights from the incorporating time series techniques .we will reiterate the impact of these time series techniques .These techniques provide valuable insights into patterns, seasonality, and trends within the data, enabling more accurate predictions and forecasts.
- Future work: It has been demonstrated in this comparative study that different time-series analysis models are capable of accurately predicting order volume in three countries Such work should include additional examinations of the impact of external factors on order volume prediction, such as gasoline prices, labour costs, service demand, and geopolitical events.

